

ECE449

Lab 3


October 4th, 2023

Lab 2




Convolutional Neural Networks (CNN)

- Build invariance to certain variations shifting, illumination, ...
- Deal with very high dimensional inputs without the need of large number of parameters
- Instead of learning to detect the whole image, detect smaller fragments in the image.

Lab 2




A set of curves




Images are made of three basic curves

- Understanding that these two images are two different instances (i.e. 'X') of the same concept is difficult for a simple neural network.




Lab 2

- A parameter matrix is called a filter.
- We need three filters to detect three curves.



Lab 2

- Align the filter and the image patch.
- Multiply each image pixel by the corresponding filter pixel.
- Add the result.
- Divide by the size of the filter.
(Implementation differs!)



Lab 2

•Activation Map

-1	-1	-1	-1	-1	-1	-1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	-1	-1	1	-1	-1	-1	-1
-1	-1	-1	1	-1	1	-1	-1	-1
-1	-1	1	-1	-1	-1	1	-1	-1
-1	1	-1	-1	-1	-1	-1	1	-1
-1	-1	-1	-1	-1	-1	-1	-1	-1

$$\begin{array}{|c|c|c|} \hline 1 & -1 & -1 \\ \hline -1 & 1 & -1 \\ \hline -1 & -1 & 1 \\ \hline \end{array} =$$

A commonly used
symbol for convolution

Activation map

0.77	-0.11	0.11	0.33	0.55	-0.11	0.33
-0.11	1.00	-0.11	0.33	-0.11	0.11	-0.11
0.11	-0.11	1.00	-0.33	0.11	-0.11	0.55
0.33	0.33	-0.33	0.55	-0.33	0.33	0.33
0.55	-0.11	0.11	-0.33	1.00	-0.11	0.11
-0.11	0.11	-0.11	0.33	-0.11	1.00	-0.11
0.33	-0.11	0.55	0.33	0.11	-0.11	0.77

Lab 2

$$\begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 \\ -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \end{bmatrix}$$

*

$$\begin{bmatrix} 1 & -1 & -1 \\ -1 & 1 & -1 \\ -1 & -1 & 1 \end{bmatrix}$$

=

$$\begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 \\ -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \end{bmatrix}$$

*

$$\begin{bmatrix} 1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & 1 \end{bmatrix}$$

=

$$\begin{bmatrix} -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & 1 & -1 & -1 & -1 & -1 \\ -1 & -1 & -1 & 1 & -1 & 1 & -1 & -1 \\ -1 & -1 & 1 & -1 & -1 & -1 & 1 & -1 \\ -1 & 1 & -1 & -1 & -1 & -1 & 1 & -1 \\ -1 & -1 & -1 & -1 & -1 & -1 & -1 & -1 \end{bmatrix}$$

*

$$\begin{bmatrix} -1 & -1 & 1 \\ -1 & 1 & -1 \\ 1 & -1 & -1 \end{bmatrix}$$

=

Forms a layer

$$\begin{bmatrix} 0.77 & -0.11 & 0.11 & 0.33 & 0.55 & -0.11 & 0.33 \\ -0.11 & 1.00 & -0.11 & 0.33 & -0.11 & 0.11 & -0.11 \\ 0.11 & -0.11 & 1.00 & -0.33 & 0.11 & -0.11 & 0.55 \\ 0.33 & 0.33 & -0.33 & 0.55 & -0.33 & 0.33 & 0.33 \\ 0.55 & -0.11 & 0.11 & -0.33 & 1.00 & -0.11 & 0.11 \\ -0.11 & 0.11 & -0.11 & 0.33 & -0.11 & 1.00 & -0.11 \\ 0.33 & -0.11 & 0.55 & 0.33 & 0.11 & -0.11 & 0.77 \end{bmatrix}$$


$$\begin{bmatrix} 0.33 & -0.55 & 0.11 & -0.11 & 0.11 & -0.55 & 0.33 \\ -0.55 & 0.55 & -0.55 & 0.33 & -0.55 & 0.55 & -0.55 \\ 0.11 & -0.55 & 0.55 & -0.77 & 0.55 & -0.55 & 0.11 \\ -0.11 & 0.33 & -0.77 & 1.00 & -0.77 & 0.33 & -0.11 \\ 0.11 & -0.55 & 0.55 & -0.77 & 0.55 & -0.55 & 0.11 \\ -0.55 & 0.55 & -0.55 & 0.33 & -0.55 & 0.55 & -0.55 \\ 0.33 & -0.55 & 0.11 & -0.11 & 0.11 & -0.55 & 0.33 \end{bmatrix}$$

$$\begin{bmatrix} 0.33 & -0.11 & 0.55 & 0.33 & 0.11 & -0.11 & 0.77 \\ -0.11 & 0.11 & -0.11 & 0.33 & -0.11 & 1.00 & -0.11 \\ 0.55 & -0.11 & 0.11 & -0.33 & 1.00 & -0.11 & 0.11 \\ 0.33 & 0.33 & -0.33 & 0.55 & -0.33 & 0.33 & 0.33 \\ 0.11 & -0.11 & 1.00 & -0.33 & 0.11 & -0.11 & 0.55 \\ -0.11 & 1.00 & -0.11 & 0.33 & -0.11 & 0.11 & -0.11 \\ 0.77 & -0.11 & 0.11 & 0.33 & 0.55 & -0.11 & 0.33 \end{bmatrix}$$


Considered
different channels
of a same
activation map

Lab 2


- Stride is the distance between two consecutive image patches used during convolution.
- Stride is measured in terms of pixels.
- If $\text{stride} < \text{patch-width}$, patches overlap.
- If $\text{stride} \geq \text{patch-width}$, patches do not overlap.



Stride = 1



Stride = 2



Stride = 3

Lab 2


Pad additional pixels on the boundary of the input to achieve a specific output size.

Ways of padding:


- Pad pixels with value equal to zero
- Repeat the boarder pixels
- Reflect the image around the boarder

Lab 2


Padding



Valid or no padding
output size $<$ input size



Same or half padding
output size $=$ input size




Full padding
output size $>$ input size

Lab 2


- The extent of the local region connected to one hidden unit is known as receptive field.
- To increase the receptive field without increasing the number of parameters, use dilation.

Dilation = 2



Lab 2

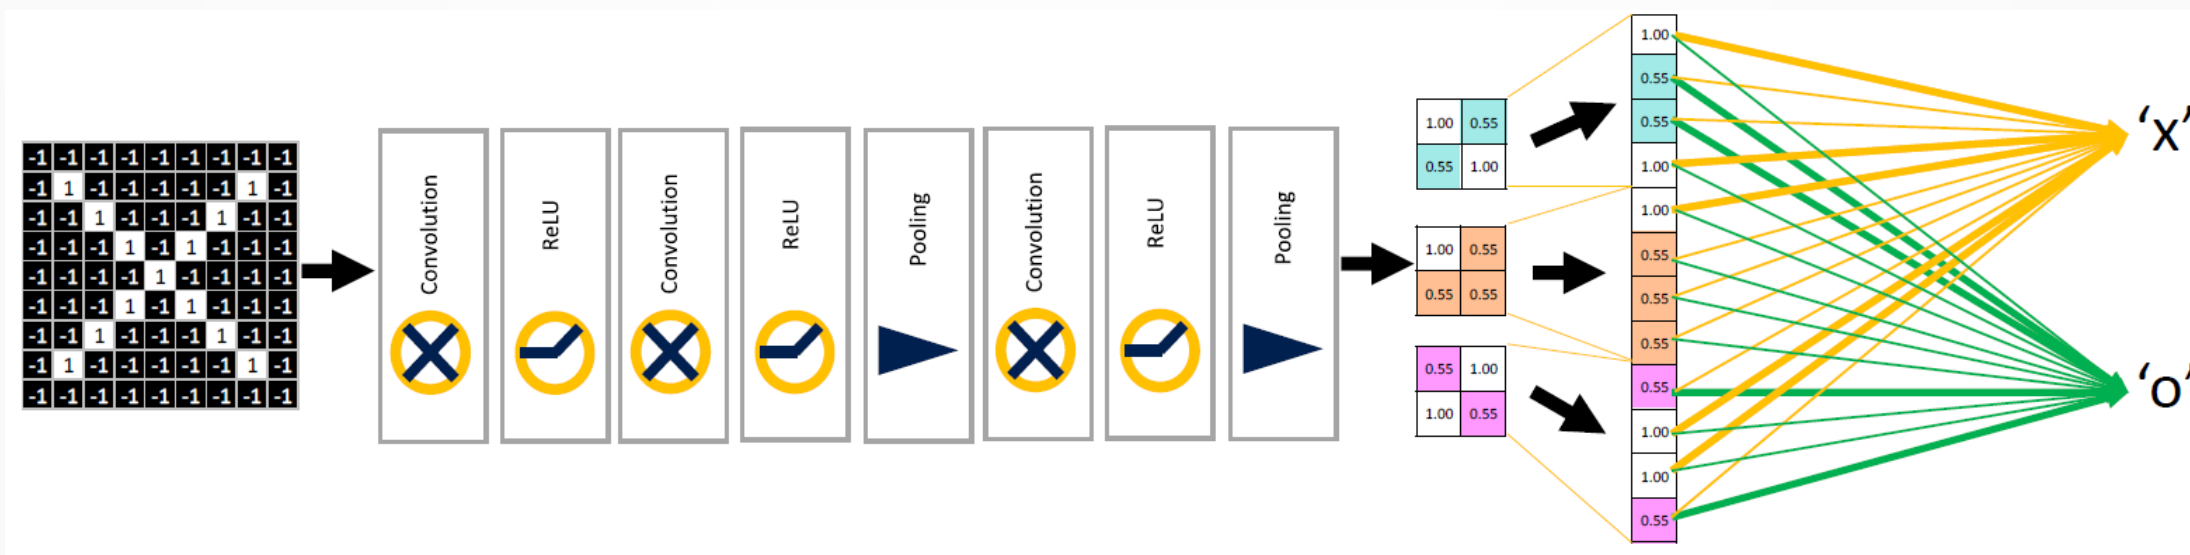
- Pool hidden units in a non overlapping neighborhood in a single channel.
- Pooling functions: max, min, mean, weighted average, etc.
- Pooling introduces invariance to local translations.



Example: max pooling

Lab 2

- In summary



Lab 2

```
from tensorflow.keras import layers, models
model = models.Sequential()
#define filters and convolutional layers here
model.add(layers.Conv2D(filters=16, kernel_size=(3, 3), activation='relu',
input_shape=(28, 28, 1)))
#Add a maxpooling layer
model.add(layers.MaxPooling2D(pool_size=(2, 2)))
#Flatten the output and give it to a fully connected layer
model.add(layers.Flatten())
#one hidden layer maps the flattened neurons to output
model.add(layers.Dense(10, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy',
metrics=['accuracy'])
```

Lab 2

- MNIST dataset: a large database of handwritten digits that is commonly used for training various image processing systems.
- Images are 28*28 pixels

```
from tensorflow.keras.datasets import mnist  
(X_train, y_train), (X_test, y_test) = mnist.load_data()
```



Lab 2

- Convert the class labels into a one-hot encoding.

- I suggest using the following:

```
from tensorflow.keras.utils import to_categorical
```

```
y_train_encoded = to_categorical(y_train)
```

```
y_test_encoded = to_categorical(y_test)
```


Lab 2

- You will solve an MNIST classification task using CNN.
- Use a ReLU activation instead of a sigmoid one.
- In your pooling layers, use the MAX() function instead of the arithmetic mean.
- For classification, use a single dense layer followed by a softmax layer.

Lab 2

```
model = models.Sequential()  
#Define filters and convolutional layers here  
model.add(layers.Conv2D(filters=16, kernel_size=(3, 3), activation='relu',  
input_shape=(28, 28, 1)))  
#Add a maxpooling layer  
model.add(layers.MaxPooling2D(pool_size=(2, 2)))  
#Flatten the output and give it to a fully connected layer  
model.add(layers.Flatten())  
#One hidden layer maps the flattened neurons to output  
model.add(layers.Dense(10, activation='softmax'))  
model.compile(optimizer=keras.optimizers.Adam(learning_rate=learning_  
rate), loss='categorical_crossentropy', metrics=['accuracy'])
```

Lab 2

- Use stratified 5-fold cross validation to test the performance of the models for at least the following parameters:
 - Number of Filters=[16, 32]
 - Learning rate = [0.001, 0.01]
- Once you have determined the best design, train the CNN one more time on the entire training set.
- Report your performance on the out-of-sample test set.

Lab 2

- You have 60,000 training images with 28×28 features, so your training will take a lot of time.